

# Homework #2 BFS Final Report

Name: Tianyu Yang

GW ID: G38878678

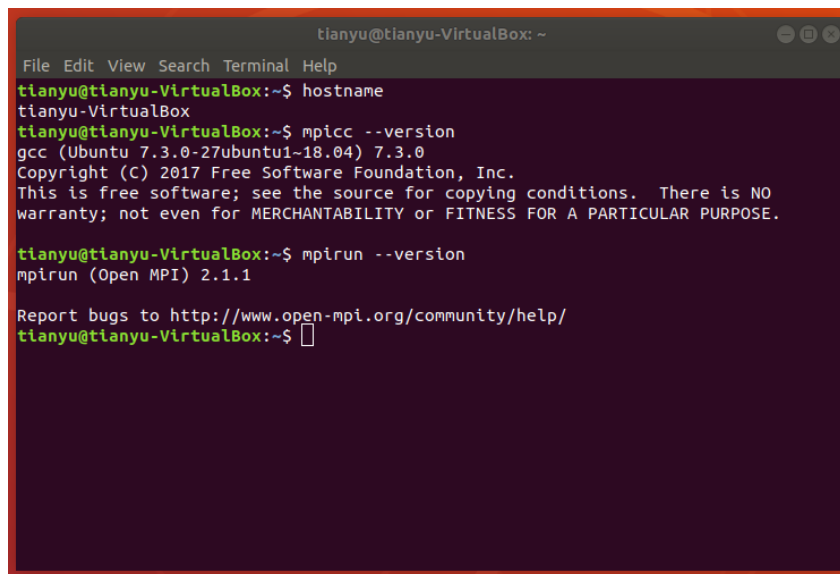
Date: 1/29/2019

## Objective:

1. include results for graphs with 100, 1K, and 1M edges. Ideally your implementation should at least outperform the reference code by 2x at this point.
2. Must include results for graphs with 1B edges. The target improvement is at least 10x over the reference code.

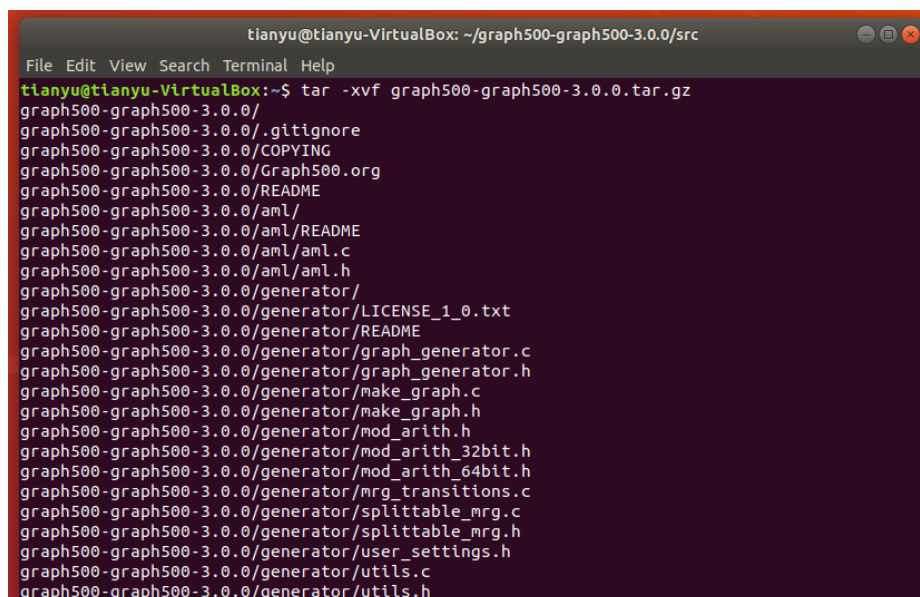
## Part 1: Preparation:

1. Install Linux on my personal workstation, using Oracle VM VirtualBox
2. Install MPI implementation on my Ubuntu installation

A terminal window titled 'tiany@tiany-VirtualBox: ~' with a menu bar (File, Edit, View, Search, Terminal, Help). The user runs 'hostname', 'mpicc --version', and 'mpirun --version'. The output shows the system is 'tiany-VirtualBox', 'mpicc' is version 7.3.0 from Ubuntu, and 'mpirun' is version 2.1.1 from Open MPI. A link to report bugs is also displayed.

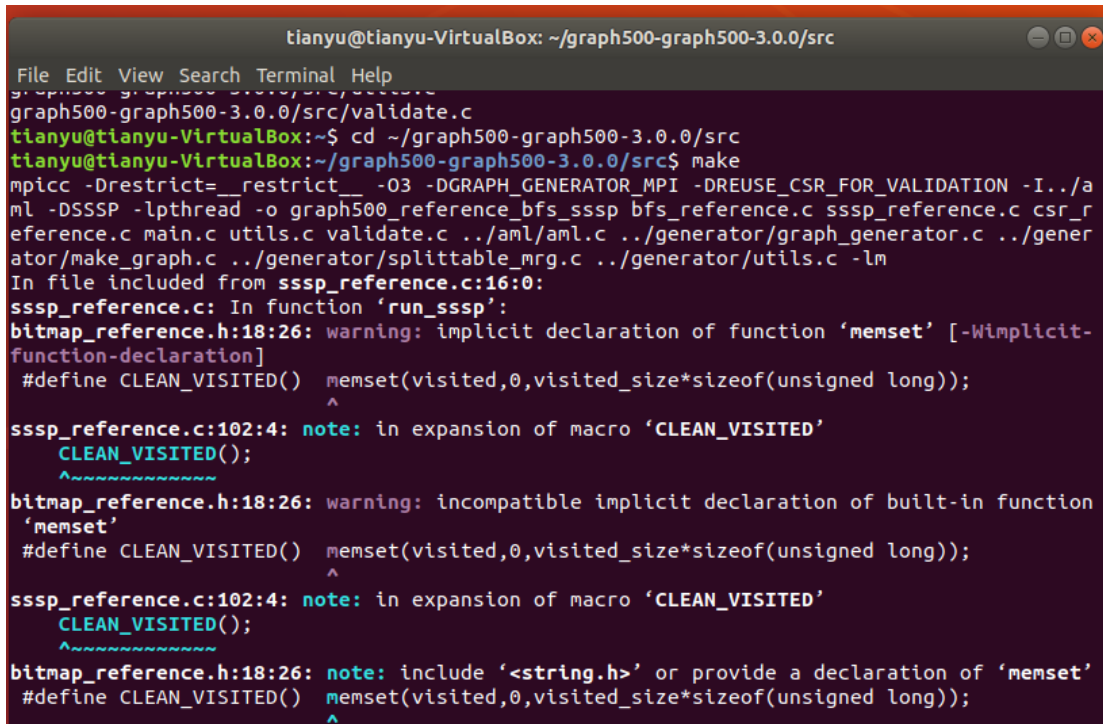
```
tiany@tiany-VirtualBox: ~  
File Edit View Search Terminal Help  
tiany@tiany-VirtualBox:~$ hostname  
tiany-VirtualBox  
tiany@tiany-VirtualBox:~$ mpicc --version  
gcc (Ubuntu 7.3.0-27ubuntu1~18.04) 7.3.0  
Copyright (C) 2017 Free Software Foundation, Inc.  
This is free software; see the source for copying conditions. There is NO  
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.  
  
tiany@tiany-VirtualBox:~$ mpirun --version  
mpirun (Open MPI) 2.1.1  
  
Report bugs to http://www.open-mpi.org/community/help/  
tiany@tiany-VirtualBox:~$
```

3. Download the reference code on Graph 500 website and decompress the tar.gz file to Home folder

A terminal window titled 'tiany@tiany-VirtualBox: ~/graph500-graph500-3.0.0/src' with a menu bar (File, Edit, View, Search, Terminal, Help). The user runs 'tar -xvf graph500-graph500-3.0.0.tar.gz'. The output lists the contents of the tar archive, including directories like 'graph500-graph500-3.0.0/' and files like 'graph500-graph500-3.0.0/.gitignore', 'COPYING', 'Graph500.org', 'README', 'aml/', 'aml/README', 'aml/aml.c', 'aml/aml.h', 'generator/', 'generator/LICENSE\_1\_0.txt', 'generator/README', 'generator/graph\_generator.c', 'generator/graph\_generator.h', 'generator/make\_graph.c', 'generator/make\_graph.h', 'generator/mod\_arith.h', 'generator/mod\_arith\_32bit.h', 'generator/mod\_arith\_64bit.h', 'generator/mrg\_transitions.c', 'generator/splittable\_mrg.c', 'generator/splittable\_mrg.h', 'generator/user\_settings.h', 'generator/utls.c', and 'generator/utls.h'.

```
tiany@tiany-VirtualBox: ~/graph500-graph500-3.0.0/src  
File Edit View Search Terminal Help  
tiany@tiany-VirtualBox:~$ tar -xvf graph500-graph500-3.0.0.tar.gz  
graph500-graph500-3.0.0/  
graph500-graph500-3.0.0/.gitignore  
graph500-graph500-3.0.0/COPYING  
graph500-graph500-3.0.0/Graph500.org  
graph500-graph500-3.0.0/README  
graph500-graph500-3.0.0/aml/  
graph500-graph500-3.0.0/aml/README  
graph500-graph500-3.0.0/aml/aml.c  
graph500-graph500-3.0.0/aml/aml.h  
graph500-graph500-3.0.0/generator/  
graph500-graph500-3.0.0/generator/LICENSE_1_0.txt  
graph500-graph500-3.0.0/generator/README  
graph500-graph500-3.0.0/generator/graph_generator.c  
graph500-graph500-3.0.0/generator/graph_generator.h  
graph500-graph500-3.0.0/generator/make_graph.c  
graph500-graph500-3.0.0/generator/make_graph.h  
graph500-graph500-3.0.0/generator/mod_arith.h  
graph500-graph500-3.0.0/generator/mod_arith_32bit.h  
graph500-graph500-3.0.0/generator/mod_arith_64bit.h  
graph500-graph500-3.0.0/generator/mrg_transitions.c  
graph500-graph500-3.0.0/generator/splittable_mrg.c  
graph500-graph500-3.0.0/generator/splittable_mrg.h  
graph500-graph500-3.0.0/generator/user_settings.h  
graph500-graph500-3.0.0/generator/utls.c  
graph500-graph500-3.0.0/generator/utls.h
```

4. Use cd command in the terminal to get into src folder and use make command to build the code



```
tianyu@tianyu-VirtualBox: ~/graph500-graph500-3.0.0/src
File Edit View Search Terminal Help
graph500-graph500-3.0.0/src/validate.c
graph500-graph500-3.0.0/src/validate.c
tianyu@tianyu-VirtualBox:~$ cd ~/graph500-graph500-3.0.0/src
tianyu@tianyu-VirtualBox:~/graph500-graph500-3.0.0/src$ make
mpicc -Drestrict=__restrict__ -O3 -DGRAPH_GENERATOR_MPI -DREUSE_CSR_FOR_VALIDATION -I../a
ml -DSSSP -lpthread -o graph500_reference_bfs_sssp bfs_reference.c sssp_reference.c csr_r
eference.c main.c utils.c validate.c ../aml/aml.c ../generator/graph_generator.c ../gener
ator/make_graph.c ../generator/splittable_mrg.c ../generator/utils.c -lm
In file included from sssp_reference.c:16:0:
sssp_reference.c: In function 'run_sssp':
bitmap_reference.h:18:26: warning: implicit declaration of function 'memset' [-Wimplicit-
function-declaration]
#define CLEAN_VISITED() memset(visited,0,visited_size*sizeof(unsigned long));
                        ^
sssp_reference.c:102:4: note: in expansion of macro 'CLEAN_VISITED'
    CLEAN_VISITED();
    ^
bitmap_reference.h:18:26: warning: incompatible implicit declaration of built-in function
'memset'
#define CLEAN_VISITED() memset(visited,0,visited_size*sizeof(unsigned long));
                        ^
sssp_reference.c:102:4: note: in expansion of macro 'CLEAN_VISITED'
    CLEAN_VISITED();
    ^
bitmap_reference.h:18:26: note: include '<string.h>' or provide a declaration of 'memset'
#define CLEAN_VISITED() memset(visited,0,visited_size*sizeof(unsigned long));
                        ^
```

5. Obtain the shared library file for reference\_bfs file and run it in the terminal
6. Adjust the edge value such as 100, 1K, 1M and 1B, and get the results for output performance such as processing time.

As for the number of edges.  $M = \text{edgefactor} * N = \text{edgefactor} * 2^{\text{SCALE}}$ .

Therefore, when the edge value is 100, edgefactor = 16, SCALE = 3; when the edge value is 1K, edgefactor = 16, SCALE = 6; when the edge is 1M, edgefactor = 16, SCALE = 16; when the edge is 1B, edgefactor = 16, SCALE = 26;

7. Compare the results and read bfs\_reference.c code

When the edge value is 100:

```
tianyu@tianyu-VirtualBox: ~/graph500-graph500-3.0.0/src
File Edit View Search Terminal Help
SCALE: 3
edgefactor: 16
NBFS: 64
graph_generation: 0.000435171
num_mpi_processes: 1
construction_time: 0.000101671
bfs_min_time: 5.1069e-05
bfs_firstquartile_time: 6.0191e-05
bfs_median_time: 6.04215e-05
bfs_thirdquartile_time: 6.08145e-05
bfs_max_time: 0.000193544
bfs_mean_time: 6.38315e-05
bfs_stddev_time: 1.95768e-05
min_nedge: 100
firstquartile_nedge: 100
median_nedge: 100
thirdquartile_nedge: 100
max_nedge: 100
mean_nedge: 100
stddev_nedge: 0
bfs_min_TEPS: 516678
bfs_firstquartile_TEPS: 1.64434e+06
bfs_median_TEPS: 1.65504e+06
bfs_thirdquartile_TEPS: 1.66138e+06
bfs_max_TEPS: 1.95814e+06
bfs_harmonic_mean_TEPS: ! 1.56663e+06
bfs_harmonic_stddev_TEPS: 60534.4
bfs_min_validate: 6.4831e-05
bfs_firstquartile_validate: 7.51395e-05
bfs_median_validate: 7.5605e-05
bfs_thirdquartile_validate: 7.85705e-05
bfs_max_validate: 0.00601966
bfs_mean_validate: 0.000235278
bfs_stddev_validate: 0.000873728
```

When the edge value is 1K:

```
tianyu@tianyu-VirtualBox: ~/graph500-graph500-3.0.0/src
File Edit View Search Terminal Help
SCALE: 6
edgefactor: 16
NBFS: 64
graph_generation: 0.000275946
num_mpi_processes: 1
construction_time: 0.000171534
bfs_min_time: 9.42e-05
bfs_firstquartile_time: 0.000106172
bfs_median_time: 0.000108293
bfs_thirdquartile_time: 0.000119494
bfs_max_time: 0.000304833
bfs_mean_time: 0.000123148
bfs_stddev_time: 3.7532e-05
min_nedge: 963
firstquartile_nedge: 963
median_nedge: 963
thirdquartile_nedge: 963
max_nedge: 963
mean_nedge: 963
stddev_nedge: 0
bfs_min_TEPS: 3.15911e+06
bfs_firstquartile_TEPS: 8.05898e+06
bfs_median_TEPS: 8.89254e+06
bfs_thirdquartile_TEPS: 9.07015e+06
bfs_max_TEPS: 1.02229e+07
bfs_harmonic_mean_TEPS: ! 7.81983e+06
bfs_harmonic_stddev_TEPS: 300262
bfs_min_validate: 0.000159405
bfs_firstquartile_validate: 0.000172178
bfs_median_validate: 0.000192394
bfs_thirdquartile_validate: 0.000239662
bfs_max_validate: 0.00711093
bfs_mean_validate: 0.000503457
bfs_stddev_validate: 0.00122803
```

When the edge value is 1M:

```
tianyu@tianyu-VirtualBox: ~/graph500-graph500-3.0.0/src
File Edit View Search Terminal Help
SCALE: 16
edgefactor: 16
NBFS: 64
graph_generation: 0.40025
num_mpi_processes: 1
construction_time: 0.155881
bfs_min_time: 0.0854266
bfs_firstquartile_time: 0.0993126
bfs_median_time: 0.101015
bfs_thirdquartile_time: 0.103822
bfs_max_time: 0.123599
bfs_mean_time: 0.101732
bfs_stddev_time: 0.0061762
min_nedge: 1048079
firstquartile_nedge: 1048079
median_nedge: 1048079
thirdquartile_nedge: 1048079
max_nedge: 1048079
mean_nedge: 1048079
stddev_nedge: 0
bfs_min_TEPS: 8.47967e+06
bfs_firstquartile_TEPS: 1.00949e+07
bfs_median_TEPS: 1.03755e+07
bfs_thirdquartile_TEPS: 1.05533e+07
bfs_max_TEPS: 1.22688e+07
bfs_harmonic_mean_TEPS: ! 1.03024e+07
bfs_harmonic_stddev_TEPS: 78801.2
bfs_min_validate: 0.148477
bfs_firstquartile_validate: 0.175853
bfs_median_validate: 0.191748
bfs_thirdquartile_validate: 0.206138
bfs_max_validate: 0.254605
bfs_mean_validate: 0.191855
bfs_stddev_validate: 0.021971
```

When the edge value is 1B: (As for 1B edges are too large for a personal computer.

Therefore, I use AWS (Amazon Web Services) as platform to run cloud computing code for BFS. I use EC2 (Amazon Elastic Compute cloud) virtual machine and also develop on Linux platform.)

```
File Edit View Search Terminal Help
SCALE: 26
edgefactor: 16
NBFS: 64
graph_generation: 531.6
num_mpi_processes: 1
construction_time: 125.199
bfs_min_time: 67.3316
bfs_firstquartile_time: 81.1544
bfs_median_time: 82.4981
bfs_thirdquartile_time: 85.8235
bfs_max_time: 99.873
bfs_mean_time: 83.9374
bfs_stddev_time: 0.00692806
min_nedge: 1048079
firstquartile_nedge: 1048079
median_nedge: 1048079
thirdquartile_nedge: 1048079
max_nedge: 1048079
mean_nedge: 1048079
stddev_nedge: 0
bfs_min_TEPS: 8.40579e+06
bfs_firstquartile_TEPS: 9.78184e+06
bfs_median_TEPS: 1.01761e+07
bfs_thirdquartile_TEPS: 1.03446e+07
bfs_max_TEPS: 1.24683e+07
bfs_harmonic_mean_TEPS: ! 1.00016e+07
bfs_harmonic_stddev_TEPS: 83308.6
bfs_min_validate: 0.168149
bfs_firstquartile_validate: 0.195301
bfs_median_validate: 0.199806
bfs_thirdquartile_validate: 0.219213
bfs_max_validate: 0.284528
bfs_mean_validate: 0.207619
bfs_stddev_validate: 0.0208945
```

Most related code in bfs\_reference.c

```
void make_graph_data_structure(const tuple_graph* const tg) {
    int i,j,k;
    convert_graph_to_oned_csr(tg, &g);
    column=g.column;
    rowstarts=g.rowstarts;

    visited_size = (g.nlocalverts + ulong_bits - 1) / ulong_bits;
    aml_register_handler(visithndl,1);
    q1 = xmalloc(g.nlocalverts*sizeof(int)); //100% of vertexes
    q2 = xmalloc(g.nlocalverts*sizeof(int));
    for(i=0;i<g.nlocalverts;i++) q1[i]=0,q2[i]=0; //touch memory
    visited = xmalloc(visited_size*sizeof(unsigned long));
}

void run_bfs(int64_t root, int64_t* pred) {
    int64_t nvisited;
    long sum;
    unsigned int i,j,k,lvl=1;
    pred_glob=pred;
    aml_register_handler(visithndl,1);

    CLEAN_VISITED();

    qc=0; sum=1; q2c=0;

    nvisited=1;
    if(VERTEX_OWNER(root) == rank) {
        pred[VERTEX_LOCAL(root)]=root;
        SET_VISITED(root);
        q1[0]=VERTEX_LOCAL(root);
        qc=1;
    }

    // While there are vertices in current level
    while(sum) {
#ifdef DEBUGSTATS
        double t0=aml_time();
        nbytes_sent=0; nbytes_rcvd=0;
#endif
        //for all vertices in current level send visit AMs to all neighbours
        for(i=0;i<qc;i++)
            for(j=rowstarts[q1[i]];j<rowstarts[q1[i]+1];j++)
```

```

        send_visit(COLUMN(j),q1[i]);
    aml_barrier();

    qc=q2c;int *tmp=q1;q1=q2;q2=tmp;
    sum=qc;
    aml_long_allsum(&sum);

    nvisited+=sum;

    q2c=0;
#ifdef DEBUGSTATS
    aml_long_allsum(&nbytes_sent);
    t0-=aml_time();
    if(!my_pe()) printf (" --lvl%d : %lld(%lld,%3.2f) visited in %5.2fs,
network
aggr %5.2fGb/s\n",lvl++,sum,nvisited,((double)nvisited/(double)g.notisolated)*1
00.0,-t0,-(double)nbytes_sent*8.0/(1.e9*t0));
#endif
    }
    aml_barrier();

}

```

## Part 2: Improved Code for 100, 1K and 1M edges

1. Modify bfs\_custom.c and main.c file code and rebuild the project by using make

command in the terminal

OpenMP method some essential function code:(bfs\_custom.c)

```

unsigned int *q1,*q2;

unsigned int qc,q2c; //pointer to first free element

void make_graph_data_structure(const tuple_graph* const tg) {
    unsigned int i,j,k;
    convert_graph_to_oned_csr(tg, &g);
    column=g.column;
    rowstarts=g.rowstarts;

    visited_size = (g.nlocalverts + ulong_bits - 1) / ulong_bits;
    aml_register_handler(visithndl,1);
}

```

```

    q1 = xmalloc(g.nlocalverts*sizeof(int)); //100% of vertexes
    q2 = xmalloc(g.nlocalverts*sizeof(int));
    for(i=0;i<g.nlocalverts;++i) q1[i]=0,q2[i]=0; //touch memory
    visited = xmalloc(visited_size*sizeof(unsigned long));
}

void run_bfs(int64_t root, int64_t* pred) {
    int64_t nvisited;
    long sum;
    unsigned int i,j,k,lvl=1;
    pred_glob=pred;
    aml_register_handler(visithndl,1);

    CLEAN_VISITED();

    qc=0; sum=1; q2c=0;

    nvisited=1;
    if(VERTEX_OWNER(root) == rank) {
        pred[VERTEX_LOCAL(root)]=root;
        SET_VISITED(root);
        q1[0]=VERTEX_LOCAL(root);
        qc=1;
    }

    // While there are vertices in current level
    while(sum) {
#ifdef DEBUGSTATS
        double t0=aml_time();
        nbytes_sent=0; nbytes_rcvd=0;
#endif
        //for all vertices in current level send visit AMs to all neighbours
#pragma omp parallel for{
            for(i=0;i<qc;++i)
                for(j=rowstarts[q1[i]];j<rowstarts[q1[i]+1];++j)
                    send_visit(COLUMN(j),q1[i]);
        }

        aml_barrier();

        qc=q2c;int *tmp=q1;q1=q2;q2=tmp;
        sum=qc;
        aml_long_allsum(&sum);

        nvisited+=sum;
    }

```



```

        q2c=0;
#ifdef DEBUGSTATS
        aml_long_allsum(&nbytes_sent);
        t0-=aml_time();
        if(!my_pe()) printf (" --lvl%d : %lld(%lld,%3.2f) visited in %5.2fs,
network
aggr %5.2fGb/s\n",lvl++,sum,nvisited,((double)nvisited/((double)g.notisolated)*1
00.0,-t0,-(double)nbytes_sent*8.0/(1.e9*t0));
#endif
    }
    aml_barrier();

}

```

Add mixing MPI and OpenMP code in main.c

```

#include "mpi.h"

#include <omp.h>

    int iam = 0, np = 1;
    MPI_Init(&argc, &argv);
    #pragma omp parallel default(shared) private(iam, np)
    {
        np = omp_get_num_threads();
        iam = omp_get_thread_num();
        run_bfs(root, &pred[0]);
    }

```

2. Adjust the edge value as 100, 1K and 1M and compare the output performance

with referenced one

For the edge value of 100:

```
tianyu@tianyu-VirtualBox: ~/graph500-graph500-3.0.0/src
File Edit View Search Terminal Help
SCALE: 3
edgefactor: 16
NBFS: 64
graph_generation: 2.45913e-05
num_mpi_processes: 1
construction_time: 9.65173e-06
bfs min_time: 4.72818e-06
bfs firstquartile_time: 5.40955e-06
bfs median_time: 5.66455e-06
bfs thirdquartile_time: 6.01482e-06
bfs max_time: 1.38396e-05
bfs mean_time: 6.0081e-06
bfs stddev_time: 1.80644e-05
min_nedge: 100
firstquartile_nedge: 100
median_nedge: 100
thirdquartile_nedge: 100
max_nedge: 100
mean_nedge: 100
stddev_nedge: 0
bfs min_TEPS: 656875
bfs firstquartile_TEPS: 1.51142e+06
bfs median_TEPS: 1.60488e+06
bfs thirdquartile_TEPS: 1.68053e+06
bfs max_TEPS: 1.92271e+06
bfs harmonic_mean_TEPS: ! 1.51311e+06
bfs harmonic_stddev_TEPS: 52106.7
bfs min_validate: 6.561e-05
bfs firstquartile_validate: 7.5771e-05
bfs median_validate: 7.6964e-05
bfs thirdquartile_validate: 8.22935e-05
bfs max_validate: 0.00429818
bfs mean_validate: 0.000148914
bfs stddev_validate: 0.000527495
```

For the edge value of 1K:

```
tianyu@tianyu-VirtualBox: ~/graph500-graph500-3.0.0/src
File Edit View Search Terminal Help
SCALE: 6
edgefactor: 16
NBFS: 64
graph_generation: 2.28317e-05
num_mpi_processes: 1
construction_time: 1.43513e-05
bfs min_time: 8.24658e-06
bfs firstquartile_time: 9.09454e-06
bfs median_time: 9.34871e-06
bfs thirdquartile_time: 1.16285e-05
bfs max_time: 4.85233e-05
bfs mean_time: 1.17164e-05
bfs stddev_time: 7.82517e-05
min_nedge: 963
firstquartile_nedge: 963
median_nedge: 963
thirdquartile_nedge: 963
max_nedge: 963
mean_nedge: 963
stddev_nedge: 0
bfs min_TEPS: 1.65384e+06
bfs firstquartile_TEPS: 6.90117e+06
bfs median_TEPS: 8.58407e+06
bfs thirdquartile_TEPS: 8.82397e+06
bfs max_TEPS: 9.7313e+06
bfs harmonic_mean_TEPS: ! 6.84938e+06
bfs harmonic_stddev_TEPS: 480286
bfs min_validate: 0.000161068
bfs firstquartile_validate: 0.000172969
bfs median_validate: 0.000178995
bfs thirdquartile_validate: 0.000251402
bfs max_validate: 0.00531467
bfs mean_validate: 0.000463966
bfs stddev_validate: 0.000922154
```

For the edge value of 1M:

```
tianyu@tianyu-VirtualBox: ~/graph500-graph500-3.0.0/src
File Edit View Search Terminal Help
SCALE: 16
edgefactor: 16
NBFS: 64
graph_generation: 0.0310456
num_mpi_processes: 1
construction_time: 0.011891
bfs_min_time: 0.00697285
bfs_firstquartile_time: 0.00836624
bfs_median_time: 0.00853069
bfs_thirdquartile_time: 0.00887255
bfs_max_time: 0.0107855
bfs_mean_time: 0.00867484
bfs_stddev_time: 0.00680403
min_nedge: 1048079
firstquartile_nedge: 1048079
median_nedge: 1048079
thirdquartile_nedge: 1048079
max_nedge: 1048079
mean_nedge: 1048079
stddev_nedge: 0
bfs_min_TEPS: 8.09788e+06
bfs_firstquartile_TEPS: 9.84383e+06
bfs_median_TEPS: 1.02383e+07
bfs_thirdquartile_TEPS: 1.04396e+07
bfs_max_TEPS: 1.25257e+07
bfs_harmonic_mean_TEPS: ! 1.00682e+07
bfs_harmonic_stddev_TEPS: 82909.5
bfs_min_validate: 0.161835
bfs_firstquartile_validate: 0.190563
bfs_median_validate: 0.195216
bfs_thirdquartile_validate: 0.208567
bfs_max_validate: 0.281916
bfs_mean_validate: 0.199185
bfs_stddev_validate: 0.0187218
```

For the edge value of 1B:

```
File Edit View Search Terminal Help
SCALE: 26
edgefactor: 16
NBFS: 64
graph_generation: 48.1188
num_mpi_processes: 1
construction_time: 11.5743
bfs_min_time: 6.09278
bfs_firstquartile_time: 7.40915
bfs_median_time: 7.54831
bfs_thirdquartile_time: 7.87305
bfs_max_time: 9.17157
bfs_mean_time: 7.68457
bfs_stddev_time: 0.00683976
min_nedge: 1048079
firstquartile_nedge: 1048079
median_nedge: 1048079
thirdquartile_nedge: 1048079
max_nedge: 1048079
mean_nedge: 1048079
stddev_nedge: 0
bfs_min_TEPS: 8.32128e+06
bfs_firstquartile_TEPS: 9.69373e+06
bfs_median_TEPS: 1.01108e+07
bfs_thirdquartile_TEPS: 1.03007e+07
bfs_max_TEPS: 1.25262e+07
bfs_harmonic_mean_TEPS: ! 9.93148e+06
bfs_harmonic_stddev_TEPS: 81097
bfs_min_validate: 0.164211
bfs_firstquartile_validate: 0.193558
bfs_median_validate: 0.197581
bfs_thirdquartile_validate: 0.212492
bfs_max_validate: 0.320519
bfs_mean_validate: 0.203395
bfs_stddev_validate: 0.0255012
```

## Part 3: Output performance for the high-performance parallel BFS code

| Edge=100       | graph_generation | construction_time | bfs mean_time |
|----------------|------------------|-------------------|---------------|
| Reference code | 0.000435171      | 0.000101671       | 0.000063815   |
| Custom code    | 2.45913E-05      | 9.65173E-06       | 6.0081E-06    |
| Multiple       | 17.69613644      | 10.53396645       | 10.62149432   |
|                |                  |                   |               |
| Edge=1K        | graph_generation | construction_time | bfs mean_time |
| Reference code | 0.000275946      | 0.000171534       | 0.000123148   |
| Custom code    | 2.28317E-05      | 1.43513E-05       | 1.17164E-05   |
| Multiple       | 12.08609083      | 11.95250604       | 10.51073709   |
|                |                  |                   |               |
| Edge=1M        | graph_generation | construction_time | bfs mean_time |
| Reference code | 0.40025          | 0.155881          | 0.101732      |
| Custom code    | 0.0310456        | 0.011891          | 0.00867484    |
| Multiple       | 12.89232613      | 13.10915819       | 11.727248     |

| Edge=1B        | graph_generation | construction_time | bfs mean_time |
|----------------|------------------|-------------------|---------------|
| Reference code | 531.6            | 125.199           | 83.9374       |
| Custom code    | 48.1188          | 11.5743           | 7.68457       |
| Multiple       | 11.04765705      | 10.81698245       | 10.92284929   |

## Part 4: Result analysis

OpenMP method:

OpenMP (Open Multi-Processing) is an application programming interface that supports multi-platform shared memory multiprocessing programming in C language. It consists of a set of compiler directives, library routines and environment variable that influence run-time behavior. OPENMP is a directory of C examples which illustrate the use of the OpenMP application program interface for carrying out parallel computations in a shared memory environment. The directives allow the user to mark areas of the code, such as do, while or for loops, which are suitable for

parallel processing. The directives appear as a special kind of comment, so the program can be compiled and run in serial mode. However, the user can tell the compiler to "notice" the special directives, in which case a version of the program will be created that runs in parallel.

OpenMP includes a number of functions whose type must be declared in any program that uses them. A user program calling OpenMP must have the statement:

***#include <omp.h>***

In this project, I use OPENMP method to decrease run-time of BFS algorithm. This method is to use multi-threads to process the nodes by using CSR format (Compressed sparse row). I use three threads to operate the data and increase the efficiency by 10x at least.

## Conclusion

By using the method above like openMP (Multi-threads), I successfully decrease the run-time of BFS by at least 10x for the edges of 100, 1K, 1M and 1B.